3.a

In [13]:
```
import numpy as np
import matplotlib.pyplot as plt

def kmeans(X, k=3, max_iters=100, tol=1e-4, random_state=None):

    if random_state is not None:
        np.random.seed(random_state)

    # Step 1: Initialize centroids randomly from data points
    n_samples = X.shape[0]
    indices = np.random.choice(n_samples, k, replace=False)
    centroids = X[indices].copy()

    # Step 2: Iterate until convergence
    for iteration in range(max_iters):
        # Assign points to nearest centroid (L2 distance)
        distances = np.sqrt(np.sum((X[:, np.newaxis, :] - centroids[np.newaxis, :,
        labels = np.argmin(distances, axis=1)

        # Update centroids to mean of assigned points
        new_centroids = np.array([X[labels == i].mean(axis=0) if np.sum(labels == i
                                  else centroids[i] for i in range(k)])

        # Check convergence
        if np.sum(np.sqrt(np.sum((new_centroids - centroids) ** 2, axis=1))) < tol:
            print(f"Converged after {iteration + 1} iterations")
            centroids = new_centroids
            break

        centroids = new_centroids

    return labels, centroids


def standardize(X):
    mean = np.mean(X, axis=0)
    std = np.std(X, axis=0)
    std[std == 0] = 1
    return (X - mean) / std, mean, std
```

In [14]:
```
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image

# Load the image
img = Image.open('jellybeans.tiff')
img_array = np.array(img)

print(f"Image shape: {img_array.shape}")
print(f"Image dtype: {img_array.dtype}")

# Visualize the original image
```

```
plt.figure(figsize=(12, 4))

plt.subplot(1, 3, 1)
plt.imshow(img_array)
plt.title('Original Jellybeans Image')
plt.axis('off')

# Reshape image to 2D array of pixels (n_pixels, 3 for RGB)
pixels = img_array.reshape(-1, 3)
print(f"\nTotal pixels: {pixels.shape[0]}")

# Standardize the pixel values
pixels_standardized, mean, std = standardize(pixels.astype(float))

# Visualize color distribution in RGB space (sample for speed)
sample_size = min(5000, len(pixels))
sample_indices = np.random.choice(len(pixels), sample_size, replace=False)
sample_pixels = pixels[sample_indices]

plt.subplot(1, 3, 2)
plt.scatter(sample_pixels[:, 0], sample_pixels[:, 1],
            c=sample_pixels/255.0, s=1, alpha=0.5)
plt.xlabel('Red')
plt.ylabel('Green')
plt.title('Color Distribution (R vs G)')
plt.grid(True, alpha=0.3)

plt.subplot(1, 3, 3)
plt.scatter(sample_pixels[:, 1], sample_pixels[:, 2],
            c=sample_pixels/255.0, s=1, alpha=0.5)
plt.xlabel('Green')
plt.ylabel('Blue')
plt.title('Color Distribution (G vs B)')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Try different K values to find the best fit
K_values = [2, 3, 4, 5, 6, 7, 8]

plt.figure(figsize=(15, 10))

for idx, k in enumerate(K_values):
    print(f"\nTrying K={k}...")
    labels, centroids = kmeans(pixels_standardized, k=k, max_iters=100, random_stat

    # Convert centroids back to original RGB scale
    centroids_rgb = centroids * std + mean
    centroids_rgb = np.clip(centroids_rgb, 0, 255).astype(np.uint8)

    # Reconstruct image with cluster colors
    reconstructed = centroids_rgb[labels].reshape(img_array.shape)

    plt.subplot(3, 3, idx + 1)
    plt.imshow(reconstructed)
```

```
    plt.title(f'K={k} clusters')
    plt.axis('off')

    # Show the dominant colors
    print(f"Dominant colors (RGB): {centroids_rgb}")

plt.tight_layout()
plt.show()
```

```
Image shape: (256, 256, 3)
Image dtype: uint8

Total pixels: 65536
```

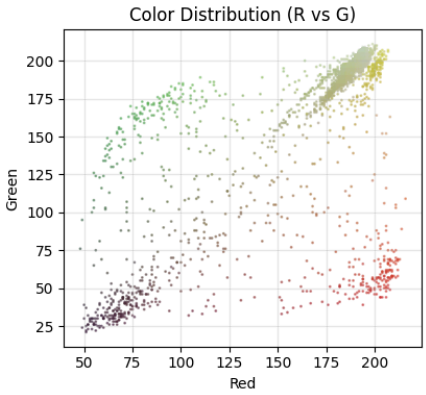

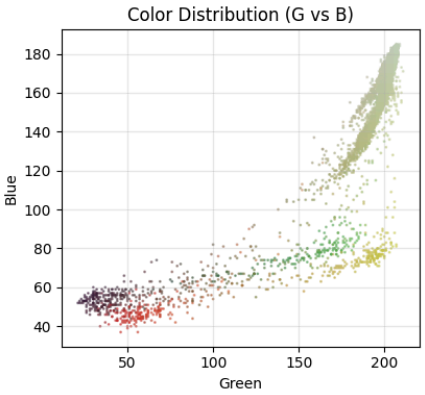

```
Trying K=2...
Converged after 8 iterations
Dominant colors (RGB): [[190 197 156]
 [118  86  60]]

Trying K=3...
Converged after 64 iterations
Dominant colors (RGB): [[189 198 161]
 [ 83  86  63]
 [185 131  67]]

Trying K=4...
Converged after 18 iterations
Dominant colors (RGB): [[191 201 167]
 [ 82  86  63]
 [181 183 117]
 [188  70  50]]

Trying K=5...
Converged after 25 iterations
Dominant colors (RGB): [[183 189 140]
 [ 81  84  62]
 [178 175  81]
 [190  65  49]
 [192 202 171]]

Trying K=6...
Converged after 55 iterations
Dominant colors (RGB): [[ 99 148  77]
 [ 75  47  55]
 [193 181  76]
 [189  65  49]
 [192 202 170]
 [182 188 138]]

Trying K=7...
Converged after 54 iterations
Dominant colors (RGB): [[ 97 148  76]
 [ 75  47  55]
 [194 181  75]
 [189  66  49]
 [193 203 174]
 [188 196 154]
 [178 182 129]]

Trying K=8...
Dominant colors (RGB): [[ 88 157  78]
 [ 72  44  54]
 [195 183  75]
 [195  62  48]
 [178 183 130]
 [189 196 154]
 [133 106  65]
 [193 203 174]]
```

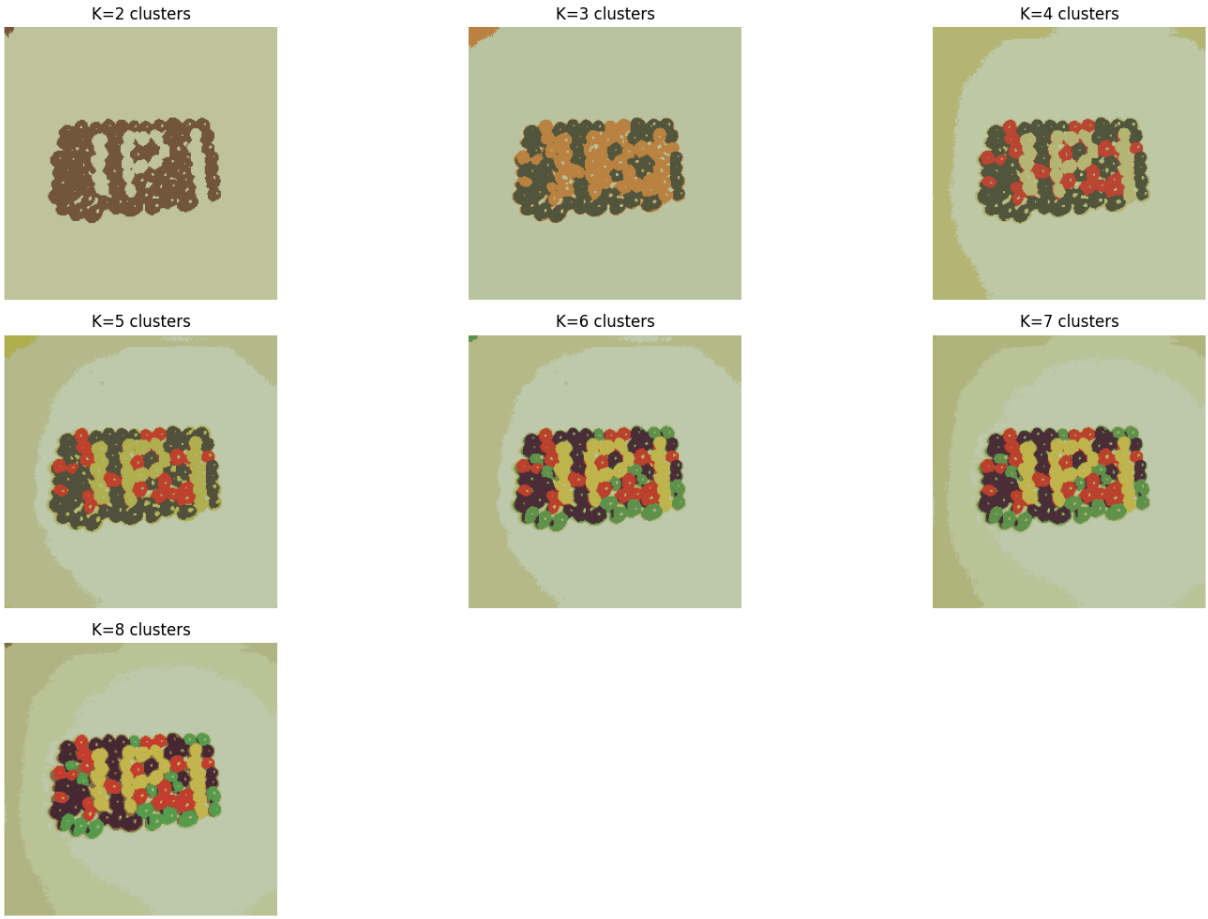

K=5 or K=6 strikes the best balance between capturing the images distinct parts while avoiding over-segmentation

3.b

In [15]:
```python
from scipy.linalg import svd
import numpy as np

def principal_component_analysis(data):
    print("Starting PCA analysis...")

    # Input validation
    if data.ndim != 2:
        raise ValueError("Data must be 2D (samples × features)")

    n_samples, n_features = data.shape
    print(f"Data shape: {data.shape} ({n_samples} samples, {n_features} features)")

    # Center data around the mean
    data_centered = data - data.mean(axis=0)

    print(f"Mean-centered data. Range: [{data_centered.min():.3f}, {data_centered.m

    # Apply SVD decomposition
    U, singular_values, Vt = svd(data_centered, full_matrices=False)

    # Compute eigenvalues from singular values
```

```
    eigenvalues = np.square(singular_values[:n_features]) / (n_samples - 1)

    # Principal components are rows of Vt transposed (columns of V)
    principal_components = Vt.T

    print(f"Eigenvalues sorted? {np.all(eigenvalues[:-1] >= eigenvalues[1:])}")
    print(f"First 5 eigenvalues: {eigenvalues[:5]}")

    return principal_components, eigenvalues, data_centered
```

In [16]:

```
import numpy as np
import matplotlib.pyplot as plt

# Load the Sentinel-2 data
print("Loading Sentinel-2 data...")
sentinel_data = np.load('sentinel2_rochester.npy')
print(f"Original data shape: {sentinel_data.shape}")
print(f"Data type: {sentinel_data.dtype}")

# The data is typically (height, width, bands)

original_shape = sentinel_data.shape
height, width, n_bands = original_shape

# Reshape to 2D: (n_pixels, n_bands)
pixels = sentinel_data.reshape(-1, n_bands)
print(f"\nReshaped to: {pixels.shape} (n_pixels, n_bands)")

# Check for any invalid values
print(f"Data range: [{pixels.min():.3f}, {pixels.max():.3f}]")
print(f"Any NaN values? {np.isnan(pixels).any()}")
print(f"Any Inf values? {np.isinf(pixels).any()}")

# Apply PCA
print("\n" + "="*60)
pcs, eigenvalues, centered_data = principal_component_analysis(pixels)
print("="*60)

# Calculate explained variance ratio
total_variance = np.sum(eigenvalues)
explained_variance_ratio = eigenvalues / total_variance
cumulative_variance = np.cumsum(explained_variance_ratio)

print(f"\nTotal variance: {total_variance:.3f}")
print(f"\nExplained variance ratio by component:")
for i in range(min(10, len(explained_variance_ratio))):
    print(f"  PC{i+1}: {explained_variance_ratio[i]:.4f} ({explained_variance_ratio

print(f"\nCumulative variance explained:")
for i in [2, 3, 4, 5]:
    if i < len(cumulative_variance):
        print(f"  First {i+1} PCs: {cumulative_variance[i]:.4f} ({cumulative_varian

# Visualize explained variance
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
plt.bar(range(1, min(11, len(eigenvalues)+1)), explained_variance_ratio[:10])
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.title('Variance Explained by Each PC')
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
plt.plot(range(1, min(11, len(cumulative_variance)+1)), cumulative_variance[:10], '
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Cumulative Variance Explained')
plt.grid(True, alpha=0.3)
plt.axhline(y=0.95, color='r', linestyle='--', label='95% threshold')
plt.legend()

plt.tight_layout()
plt.show()

# Extract different numbers of principal components
n_components_list = [3, 4, 5, 6]
K = 6

print("\n" + "="*60)
print("CLUSTERING WITH DIFFERENT NUMBERS OF PCs")
print("="*60)

# Store results
results = {}

for n_comp in n_components_list:
    print(f"\n--- Using {n_comp} Principal Components ---")

    # Project data onto first n_comp principal components
    # Projection: centered_data @ pcs[:, :n_comp]
    projected_data = centered_data @ pcs[:, :n_comp]
    print(f"Projected data shape: {projected_data.shape}")
    print(f"Projected data range: [{projected_data.min():.3f}, {projected_data.max(

    # Standardize the projected data for K-Means
    projected_standardized, proj_mean, proj_std = standardize(projected_data)

    # Apply K-Means clustering
    print(f"Applying K-Means with K={K}...")
    labels, centroids = kmeans(projected_standardized, k=K, max_iters=100, random_s

    # Reshape labels back to image dimensions
    label_image = labels.reshape(height, width)

    # Store results
    results[n_comp] = {
        'projected_data': projected_data,
        'labels': labels,
        'label_image': label_image,
        'centroids': centroids,
        'variance_explained': cumulative_variance[n_comp-1]
```

```
    }

    print(f"Clustering complete!")
    print(f"Unique clusters: {np.unique(labels)}")
    print(f"Cluster sizes: {[np.sum(labels == i) for i in range(K)]}")
    print(f"Variance explained with {n_comp} PCs: {cumulative_variance[n_comp-1]*10

# Visualize clustering results
plt.figure(figsize=(16, 10))

for idx, n_comp in enumerate(n_components_list):
    label_image = results[n_comp]['label_image']
    variance = results[n_comp]['variance_explained']

    plt.subplot(2, 2, idx + 1)
    plt.imshow(label_image, cmap='tab10')
    plt.title(f'{n_comp} PCs (K={K} clusters)\nVariance: {variance*100:.2f}%')
    plt.colorbar(label=f'Cluster ID', ticks=range(K))
    plt.axis('off')

plt.suptitle(f'K-Means Clustering on PCA-Reduced Sentinel-2 Data', fontsize=14, fon
plt.tight_layout()
plt.show()

# Visualize first 3 PCs as RGB image
print("\n" + "="*60)
print("VISUALIZING FIRST 3 PCs AS RGB IMAGE")
print("="*60)

first_3_pcs = (centered_data @ pcs[:, :3]).reshape(height, width, 3)

# Normalize to 0-1 range for visualization
pc_min = first_3_pcs.min()
pc_max = first_3_pcs.max()
first_3_pcs_normalized = (first_3_pcs - pc_min) / (pc_max - pc_min)

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.imshow(first_3_pcs_normalized)
plt.title('First 3 PCs as RGB\n(PC1=Red, PC2=Green, PC3=Blue)')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(results[3]['label_image'], cmap='tab10')
plt.title(f'K-Means Clustering (3 PCs, K={K})')
plt.colorbar(label='Cluster ID', ticks=range(K))
plt.axis('off')

plt.tight_layout()
plt.show()

# Compare cluster distributions
print("\n" + "="*60)
print("CLUSTER DISTRIBUTION COMPARISON")
print("="*60)
```

```python
for n_comp in n_components_list:
    labels = results[n_comp]['labels']
    print(f"\n{n_comp} PCs:")
    for cluster_id in range(K):
        count = np.sum(labels == cluster_id)
        percentage = (count / len(labels)) * 100
        print(f"  Cluster {cluster_id}: {count:7d} pixels ({percentage:5.2f}%)")

# Optional: Show which number of PCs gives the most distinct clustering
print("\n" + "="*60)
print("RECOMMENDATIONS")
print("="*60)
print(f"Consider the trade-off between:")
print(f"  - Fewer PCs (3-4): Faster computation, less noise, but may lose important
print(f"  - More PCs (5-6): More information retained, but potentially more noise")
print(f"\nBased on cumulative variance:")
for n_comp in n_components_list:
    var = results[n_comp]['variance_explained']
    print(f"  {n_comp} PCs capture {var*100:.2f}% of variance")
```

```
Loading Sentinel-2 data...
Original data shape: (954, 716, 12)
Data type: float64

Reshaped to: (683064, 12) (n_pixels, n_bands)
Data range: [0.000, 0.929]
Any NaN values? False
Any Inf values? False

============================================================
Starting PCA analysis...
Data shape: (683064, 12) (683064 samples, 12 features)
Mean-centered data. Range: [-0.318, 0.810]
Eigenvalues sorted? True
First 5 eigenvalues: [0.06332655 0.01003501 0.00177272 0.00085874 0.00018644]
============================================================

Total variance: 0.076

Explained variance ratio by component:
  PC1: 0.8282 (82.82%)
  PC2: 0.1312 (13.12%)
  PC3: 0.0232 (2.32%)
  PC4: 0.0112 (1.12%)
  PC5: 0.0024 (0.24%)
  PC6: 0.0010 (0.10%)
  PC7: 0.0010 (0.10%)
  PC8: 0.0007 (0.07%)
  PC9: 0.0005 (0.05%)
  PC10: 0.0004 (0.04%)

Cumulative variance explained:
  First 3 PCs: 0.9827 (98.27%)
  First 4 PCs: 0.9939 (99.39%)
  First 5 PCs: 0.9963 (99.63%)
  First 6 PCs: 0.9973 (99.73%)
```

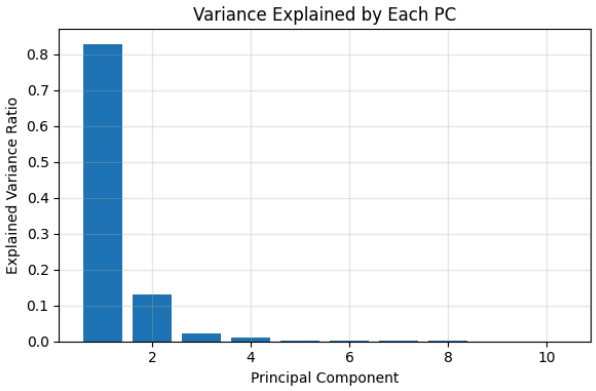

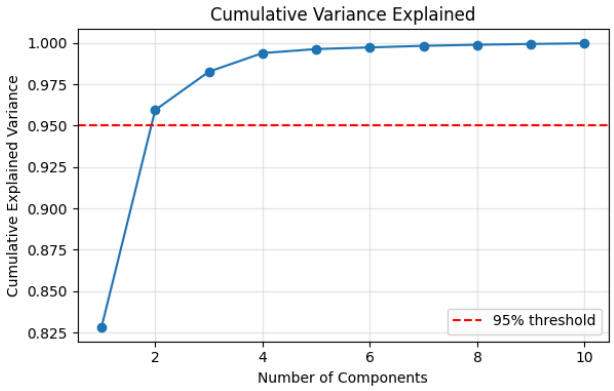

```
============================================================
CLUSTERING WITH DIFFERENT NUMBERS OF PCs
============================================================

--- Using 3 Principal Components ---
Projected data shape: (683064, 3)
Projected data range: [-1.165, 1.406]
Applying K-Means with K=6...
Converged after 88 iterations
Clustering complete!
Unique clusters: [0 1 2 3 4 5]
Cluster sizes: [np.int64(69414), np.int64(147856), np.int64(80775), np.int64(18149
1), np.int64(37890), np.int64(165638)]
Variance explained with 3 PCs: 98.27%

--- Using 4 Principal Components ---
Projected data shape: (683064, 4)
Projected data range: [-1.165, 1.406]
Applying K-Means with K=6...
Converged after 55 iterations
Clustering complete!
Unique clusters: [0 1 2 3 4 5]
Cluster sizes: [np.int64(69966), np.int64(2676), np.int64(129068), np.int64(196959),
np.int64(40638), np.int64(243757)]
Variance explained with 4 PCs: 99.39%

--- Using 5 Principal Components ---
Projected data shape: (683064, 5)
Projected data range: [-1.165, 1.406]
Applying K-Means with K=6...
Converged after 42 iterations
Clustering complete!
Unique clusters: [0 1 2 3 4 5]
Cluster sizes: [np.int64(54725), np.int64(2532), np.int64(84435), np.int64(233340),
np.int64(51469), np.int64(256563)]
Variance explained with 5 PCs: 99.63%

--- Using 6 Principal Components ---
Projected data shape: (683064, 6)
Projected data range: [-1.165, 1.406]
Applying K-Means with K=6...
Converged after 39 iterations
Clustering complete!
Unique clusters: [0 1 2 3 4 5]
Cluster sizes: [np.int64(53949), np.int64(151801), np.int64(77540), np.int64(16952
6), np.int64(38054), np.int64(192194)]
Variance explained with 6 PCs: 99.73%
```

K-Means Clustering on PCA-Reduced Sentinel-2 Data

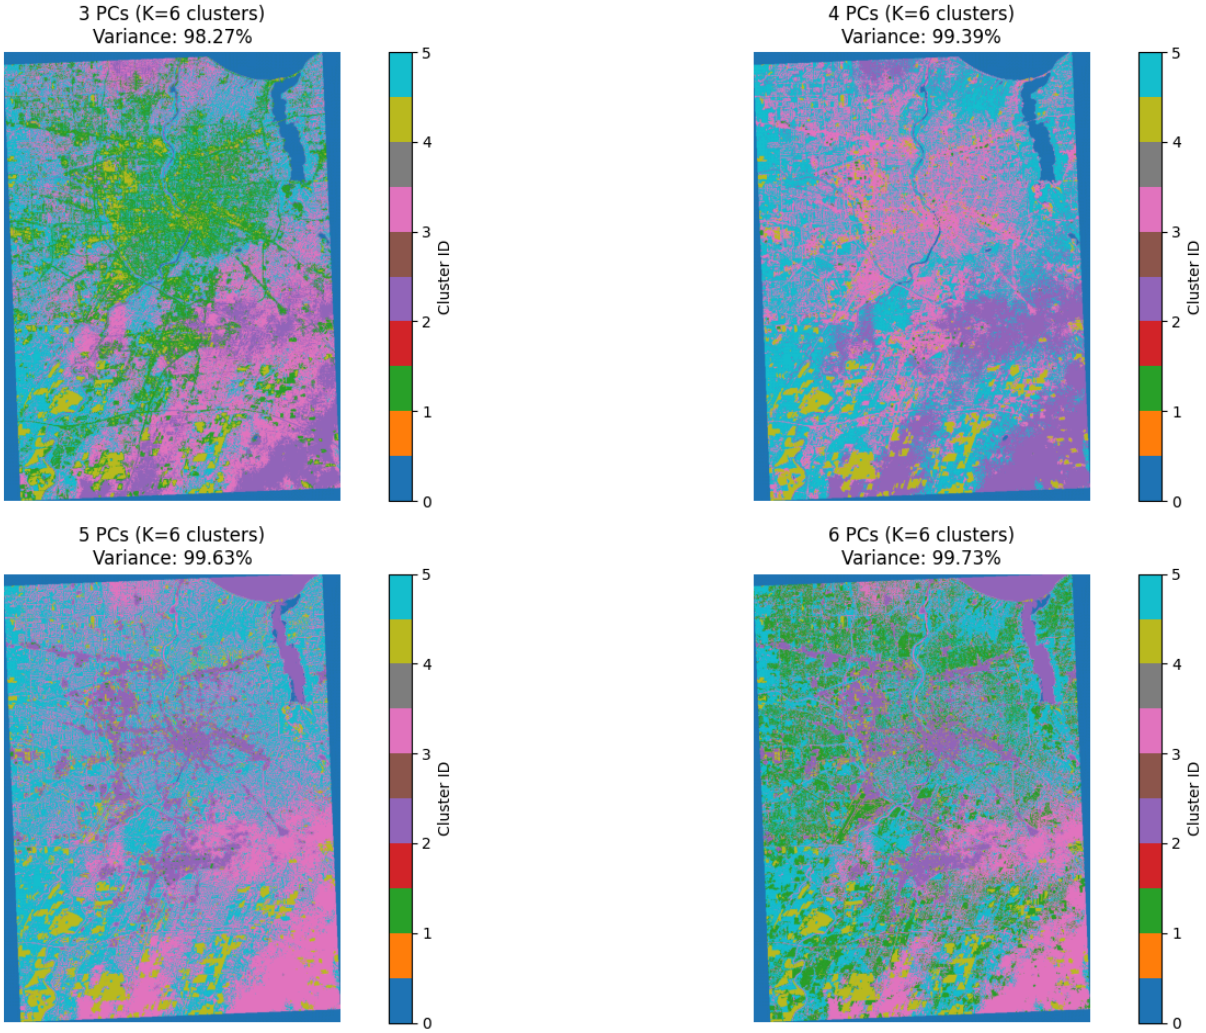


```
============================================================
VISUALIZING FIRST 3 PCs AS RGB IMAGE
============================================================
```

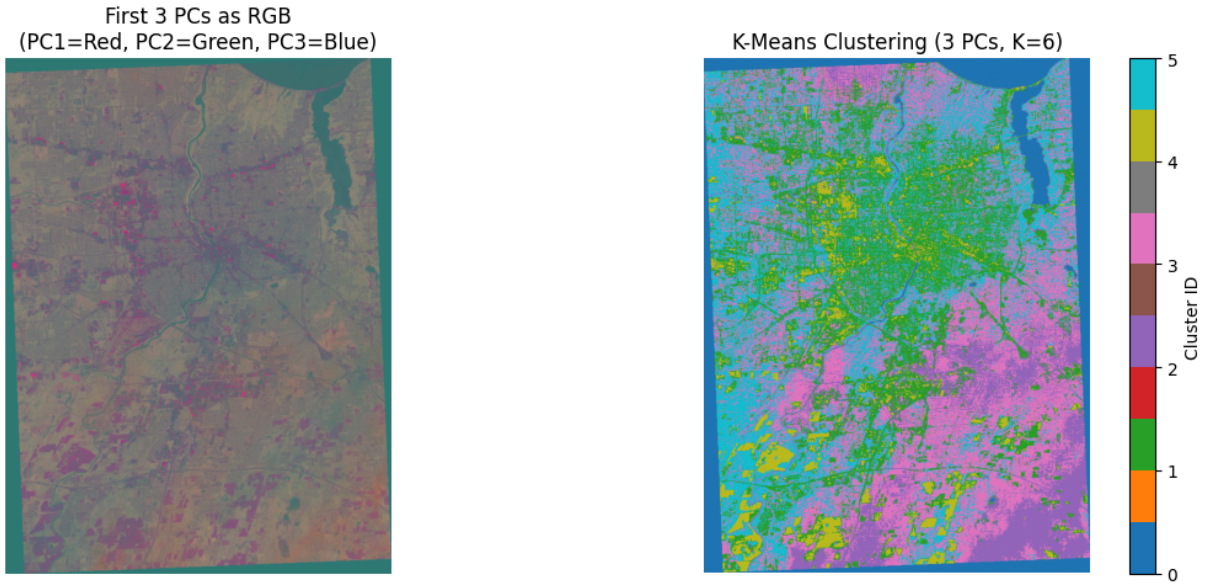

```
============================================================
CLUSTER DISTRIBUTION COMPARISON
============================================================

3 PCs:
  Cluster 0:   69414 pixels (10.16%)
  Cluster 1:  147856 pixels (21.65%)
  Cluster 2:   80775 pixels (11.83%)
  Cluster 3:  181491 pixels (26.57%)
  Cluster 4:   37890 pixels ( 5.55%)
  Cluster 5:  165638 pixels (24.25%)

4 PCs:
  Cluster 0:   69966 pixels (10.24%)
  Cluster 1:    2676 pixels ( 0.39%)
  Cluster 2:  129068 pixels (18.90%)
  Cluster 3:  196959 pixels (28.83%)
  Cluster 4:   40638 pixels ( 5.95%)
  Cluster 5:  243757 pixels (35.69%)

5 PCs:
  Cluster 0:   54725 pixels ( 8.01%)
  Cluster 1:    2532 pixels ( 0.37%)
  Cluster 2:   84435 pixels (12.36%)
  Cluster 3:  233340 pixels (34.16%)
  Cluster 4:   51469 pixels ( 7.54%)
  Cluster 5:  256563 pixels (37.56%)

6 PCs:
  Cluster 0:   53949 pixels ( 7.90%)
  Cluster 1:  151801 pixels (22.22%)
  Cluster 2:   77540 pixels (11.35%)
  Cluster 3:  169526 pixels (24.82%)
  Cluster 4:   38054 pixels ( 5.57%)
  Cluster 5:  192194 pixels (28.14%)

============================================================
RECOMMENDATIONS
============================================================
Consider the trade-off between:
  - Fewer PCs (3-4): Faster computation, less noise, but may lose important features
  - More PCs (5-6): More information retained, but potentially more noise

Based on cumulative variance:
  3 PCs capture 98.27% of variance
  4 PCs capture 99.39% of variance
  5 PCs capture 99.63% of variance
  6 PCs capture 99.73% of variance
```

Explanation: K-Means Clustering on Data Reduced by PCA: The four panels show the effects of varying the number of principal components on clustering outcomes while preserving K=6 clusters. The clustering yields distinct spatial patterns with 3 PCs (98.27% variance): cyan denotes water bodies, pink/magenta covers a large portion of the central region, and green clusters predominate in the upper-left agricultural area. The cluster boundaries get more

precise and subtle as you go up to 4, 5, and 6 PCs (capturing 99.39%, 99.63%, and 99.73% variance, respectively). Observe how the pink/purple regions differentiate more clearly and the green vegetated areas exhibit more internal structure in the 6 PC result, which displays the most detailed segmentation with superior discrimination between subtle land cover variations.

First 3 PCs as RGB Visualization The RGB composite (PC1=Red, PC2=Green, and PC3=Blue) on the left produces a false-color visualization in which various hues correspond to different combinations of the three dominant spectral patterns. Agricultural fields appear as brown/gray (balanced across PCs), water appears as bright cyan/teal (strong in PC3), and the green/pink areas represent various principal component combinations. This can be compared to the K-means clustering on the right to see how the clustering algorithm groups pixels with similar PC values. The cyan water bodies are consistently identified, while the pink areas in both images correspond to similar land cover.

3.c

In [17]:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import MiniBatchKMeans
from spectral import open_image
import time

# Load the hyperspectral data
print("="*70)
print("LOADING HYPERSPECTRAL DATA")
print("="*70)

# Load using spectral library (handles .hdr files)
img = open_image('tait_labsphere.hdr')
print(f"Image shape: {img.shape}")  # (rows, cols, bands)
print(f"Number of bands: {img.shape[2]}")
print(f"Data type: {img.dtype}")

# Load the full data into memory
hyperspectral_data = img.load()
print(f"Data loaded. Shape: {hyperspectral_data.shape}")
print(f"Data range: [{hyperspectral_data.min():.3f}, {hyperspectral_data.max():.3f}

# Visualize the full image (using RGB approximation)

height, width, n_bands = hyperspectral_data.shape

# Create RGB visualization
rgb_bands = [50, 30, 20]
if n_bands >= max(rgb_bands):
    rgb_image = hyperspectral_data[:, :, rgb_bands]
    # Normalize for display
    rgb_normalized = (rgb_image - rgb_image.min()) / (rgb_image.max() - rgb_image.m
else:
    # Use first 3 bands if not enough bands
```

```
    rgb_image = hyperspectral_data[:, :, :3]
    rgb_normalized = (rgb_image - rgb_image.min()) / (rgb_image.max() - rgb_image.m

plt.figure(figsize=(12, 10))
plt.imshow(rgb_normalized)
plt.title('Full Hyperspectral Image (RGB Approximation)\nSelect a 250×250 patch wit
plt.colorbar()
plt.grid(True, alpha=0.3)
plt.xlabel('Column')
plt.ylabel('Row')
plt.show()

# Extract a 250×250 patch

print("\n" + "="*70)
print("EXTRACTING 250×250 PATCH")
print("="*70)


start_row = max(0, (height - 250) // 2)  # Center region
start_col = max(0, (width - 250) // 2)


end_row = min(start_row + 250, height)
end_col = min(start_col + 250, width)
start_row = max(0, end_row - 250)
start_col = max(0, end_col - 250)

patch = hyperspectral_data[start_row:end_row, start_col:end_col, :]
print(f"Extracted patch from ({start_row}:{end_row}, {start_col}:{end_col})")
print(f"Patch shape: {patch.shape}")

# Visualize the extracted patch
patch_rgb = patch[:, :, rgb_bands if n_bands >= max(rgb_bands) else [0, 1, 2]]
patch_rgb_norm = (patch_rgb - patch_rgb.min()) / (patch_rgb.max() - patch_rgb.min()

plt.figure(figsize=(10, 10))
plt.imshow(patch_rgb_norm)
plt.title(f'Extracted 250×250 Patch\nLocation: Row {start_row}-{end_row}, Col {star
plt.colorbar()
plt.grid(True, alpha=0.3)
plt.show()

# Reshape patch for PCA and clustering
patch_height, patch_width, patch_bands = patch.shape
pixels = patch.reshape(-1, patch_bands)
print(f"\nReshaped patch: {pixels.shape} (n_pixels={pixels.shape[0]}, n_bands={pixe

# Check for invalid values
print(f"Any NaN values? {np.isnan(pixels).any()}")
print(f"Any Inf values? {np.isinf(pixels).any()}")

# Apply PCA
print("\n" + "="*70)
print("APPLYING PCA FOR DIMENSIONALITY REDUCTION")
print("="*70)
```

```
pcs, eigenvalues, centered_data = principal_component_analysis(pixels)

# Calculate explained variance
total_variance = np.sum(eigenvalues)
explained_variance_ratio = eigenvalues / total_variance
cumulative_variance = np.cumsum(explained_variance_ratio)

print(f"\nExplained variance by first 10 components:")
for i in range(min(10, len(explained_variance_ratio))):
    print(f"  PC{i+1}: {explained_variance_ratio[i]*100:.2f}%")

print(f"\nCumulative variance:")
for n in [2, 5, 10, 50, 100]:
    if n <= len(cumulative_variance):
        print(f"  First {n:3d} PCs: {cumulative_variance[n-1]*100:.2f}%")

# Plot variance explained
plt.figure(figsize=(14, 5))

plt.subplot(1, 2, 1)
plt.plot(range(1, min(101, len(explained_variance_ratio)+1)),
         explained_variance_ratio[:100], 'b-', linewidth=2)
plt.xlabel('Principal Component', fontsize=12)
plt.ylabel('Explained Variance Ratio', fontsize=12)
plt.title('Variance Explained by Each PC', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.xlim(0, 100)

plt.subplot(1, 2, 2)
plt.plot(range(1, min(101, len(cumulative_variance)+1)),
         cumulative_variance[:100], 'r-', linewidth=2)
plt.xlabel('Number of Components', fontsize=12)
plt.ylabel('Cumulative Explained Variance', fontsize=12)
plt.title('Cumulative Variance Explained', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.axhline(y=0.95, color='g', linestyle='--', linewidth=2, label='95% threshold')
plt.axhline(y=0.99, color='orange', linestyle='--', linewidth=2, label='99% thresho
plt.legend()
plt.xlim(0, 100)

plt.tight_layout()
plt.show()

# Determine K using elbow method on a small subset
print("\n" + "="*70)
print("DETERMINING OPTIMAL K USING ELBOW METHOD")
print("="*70)

# Use 10 PCs for K selection
projected_10pc = centered_data @ pcs[:, :10]
projected_10pc_std, _, _ = standardize(projected_10pc)

K_range = range(2, 13)
inertias = []
```

```
for k in K_range:
    print(f"Testing K={k}...", end=" ")
    kmeans_test = MiniBatchKMeans(n_clusters=k, random_state=42, batch_size=1000,
                                  max_iter=100, n_init=3)
    kmeans_test.fit(projected_10pc_std)
    inertias.append(kmeans_test.inertia_)
    print(f"Inertia: {kmeans_test.inertia_:.2f}")

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(K_range, inertias, 'bo-', linewidth=2, markersize=8)
plt.xlabel('Number of Clusters (K)', fontsize=12)
plt.ylabel('Inertia', fontsize=12)
plt.title('Elbow Method for Optimal K', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.xticks(K_range)

plt.subplot(1, 2, 2)
rate_change = -np.diff(inertias)
plt.plot(list(K_range)[1:], rate_change, 'ro-', linewidth=2, markersize=8)
plt.xlabel('Number of Clusters (K)', fontsize=12)
plt.ylabel('Reduction in Inertia', fontsize=12)
plt.title('Rate of Inertia Decrease', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.xticks(list(K_range)[1:])

plt.tight_layout()
plt.show()

# Choose K based on elbow
K = 8  # Adjust this based on the elbow curve
print(f"\nSelected K = {K} clusters")

print("\n" + "="*70)

print("="*70)

print(f"  - patch: {patch.shape}")
print(f"  - pixels: {pixels.shape}")
print(f"  - patch_height, patch_width, patch_bands: {patch_height}, {patch_width},
print(f"  - pcs: {pcs.shape}")
print(f"  - centered_data: {centered_data.shape}")
print(f"  - cumulative_variance: {cumulative_variance.shape}")
print(f"  - patch_rgb_norm: {patch_rgb_norm.shape}")
print(f"  - K: {K}")

print("="*70)
```

```
======================================================================
LOADING HYPERSPECTRAL DATA
======================================================================
Image shape: (1039, 1087, 272)
Number of bands: 272
Data type: <f4
Data loaded. Shape: (1039, 1087, 272)
Data range: [0.000, 10.668]
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_12680\4003742175.py:32: DeprecationWarni
ng: __array_wrap__ must accept context and return_scalar arguments (positionally) in
the future. (Deprecated NumPy 2.0)
  rgb_normalized = (rgb_image - rgb_image.min()) / (rgb_image.max() - rgb_image.min
())
```

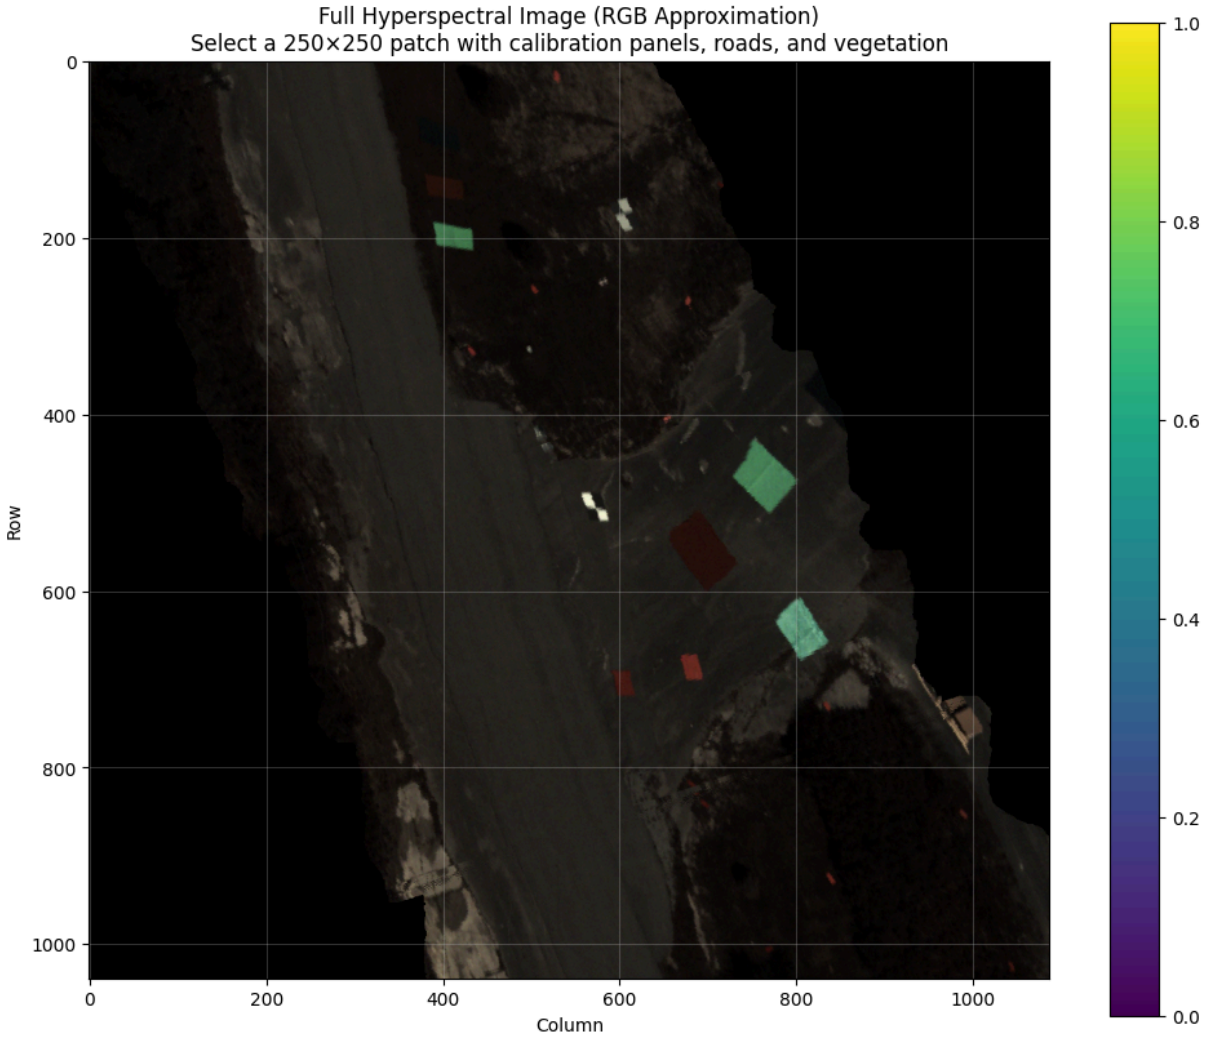


```
======================================================================
EXTRACTING 250×250 PATCH
======================================================================
Extracted patch from (394:644, 418:668)
Patch shape: (250, 250, 272)
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_12680\4003742175.py:69: DeprecationWarni
ng: __array_wrap__ must accept context and return_scalar arguments (positionally) in
the future. (Deprecated NumPy 2.0)
  patch_rgb_norm = (patch_rgb - patch_rgb.min()) / (patch_rgb.max() - patch_rgb.min
())
```

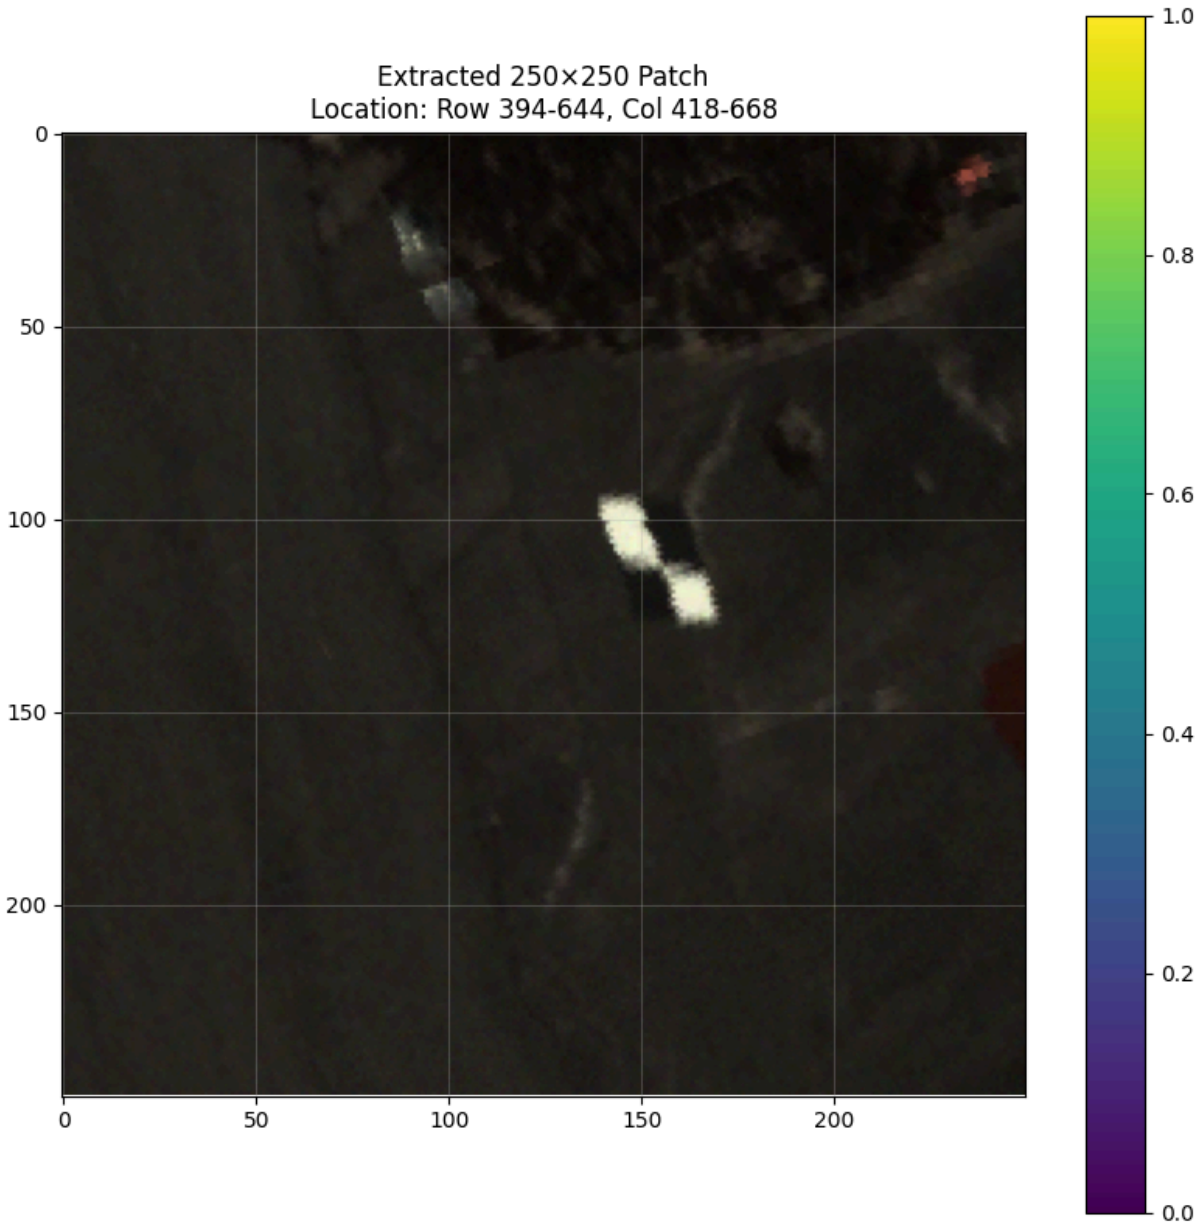
Extracted 250×250 Patch
Location: Row 394-644, Col 418-668
0
50
100
150
200
0
50
100
150
200
1.0
0.8
0.6
0.4
0.2
0.0

```
Reshaped patch: (62500, 272) (n_pixels=62500, n_bands=272)
Any NaN values? False
Any Inf values? False

======================================================================
APPLYING PCA FOR DIMENSIONALITY REDUCTION
======================================================================
Starting PCA analysis...
Data shape: (62500, 272) (62500 samples, 272 features)
Mean-centered data. Range: [-1.450, 9.355]
Eigenvalues sorted? True
First 5 eigenvalues: [32.195763   16.244665    0.6766196   0.39435345  0.1655728 ]

Explained variance by first 10 components:
  PC1: 59.97%
  PC2: 30.26%
  PC3: 1.26%
  PC4: 0.73%
  PC5: 0.31%
  PC6: 0.27%
  PC7: 0.26%
  PC8: 0.24%
  PC9: 0.22%
  PC10: 0.22%

Cumulative variance:
  First   2 PCs: 90.23%
  First   5 PCs: 92.54%
  First  10 PCs: 93.76%
  First  50 PCs: 97.05%
  First 100 PCs: 98.43%
```

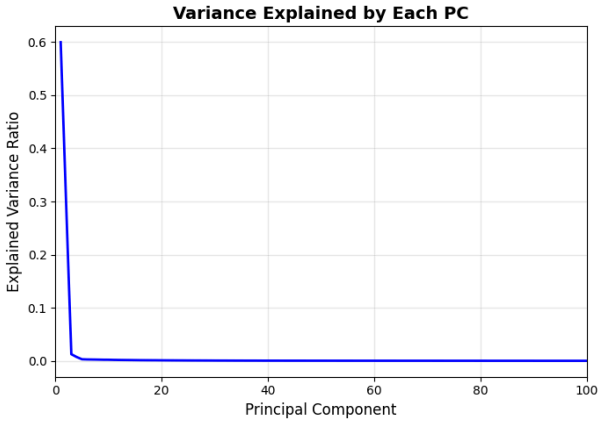

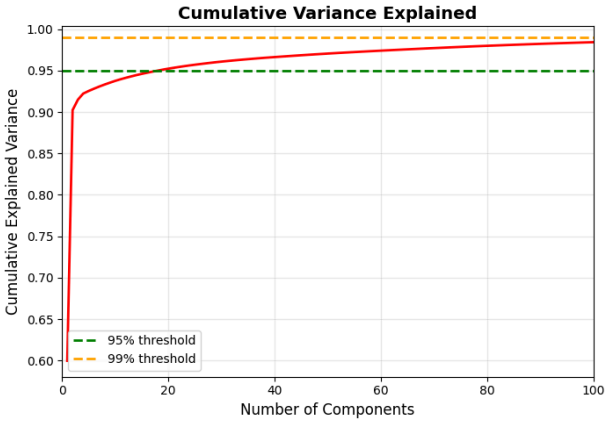

```
======================================================================
DETERMINING OPTIMAL K USING ELBOW METHOD
======================================================================
Testing K=2... Inertia: 577787.06
Testing K=3... Inertia: 542156.88
Testing K=4... Inertia: 508344.41
Testing K=5... Inertia: 492995.38
Testing K=6... Inertia: 480588.56
Testing K=7... Inertia: 449207.44
Testing K=8... Inertia: 408820.56
Testing K=9... Inertia: 420370.28
Testing K=10... Inertia: 395291.81
Testing K=11... Inertia: 381160.28
Testing K=12... Inertia: 376780.00
```

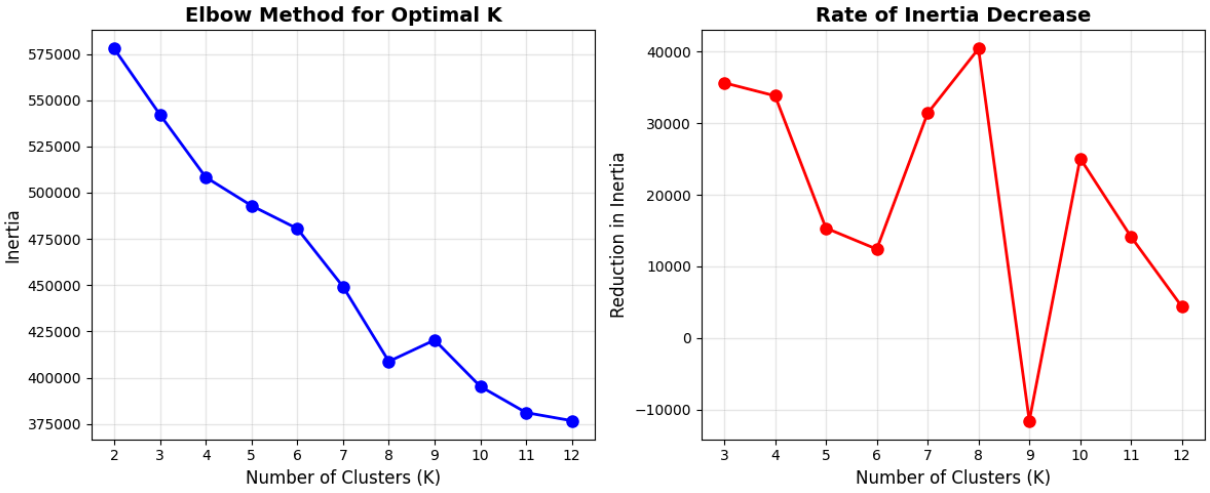


```
Selected K = 8 clusters

======================================================================
======================================================================
  - patch: (250, 250, 272)
  - pixels: (62500, 272)
  - patch_height, patch_width, patch_bands: 250, 250, 272
  - pcs: (272, 272)
  - centered_data: (62500, 272)
  - cumulative_variance: (272,)
  - patch_rgb_norm: (250, 250, 3)
  - K: 8
======================================================================
```

In [18]:

```python
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import MiniBatchKMeans
from sklearn.metrics import adjusted_rand_score
import time


print("="*70)

print("="*70)

required_vars = ['patch', 'pixels', 'patch_height', 'patch_width', 'patch_bands',
                 'pcs', 'centered_data', 'cumulative_variance', 'patch_rgb_norm', '
```

```
missing_vars = []
for var in required_vars:
    if var not in globals():
        missing_vars.append(var)

if missing_vars:
    print(f"ERROR: Missing variables: {missing_vars}")

    raise NameError(f"Missing required variables from Snippet 1: {missing_vars}")
else:
    print("✓ All required variables found!")
    print(f"  Patch shape: ({patch_height}, {patch_width}, {patch_bands})")
    print(f"  Number of pixels: {pixels.shape[0]}")
    print(f"  Number of bands: {pixels.shape[1]}")
    print(f"  K value: {K}")

# Apply MiniBatch K-Means with different numbers of PCs
print("\n" + "="*70)
print("CLUSTERING WITH DIFFERENT NUMBERS OF PRINCIPAL COMPONENTS")
print("="*70)

n_components_list = [2, 5, 10, 50, 100]
results = {}

for n_comp in n_components_list:
    print(f"\n{'='*70}")
    print(f"Processing with {n_comp} Principal Components")
    print(f"{'='*70}")

    # Project data
    projected = centered_data @ pcs[:, :n_comp]
    projected_std, _, _ = standardize(projected)

    print(f"Projected shape: {projected_std.shape}")
    print(f"Variance explained: {cumulative_variance[n_comp-1]*100:.2f}%")

    # Apply MiniBatch K-Means
    start_time = time.time()
    kmeans_model = MiniBatchKMeans(
        n_clusters=K,
        random_state=42,
        batch_size=1000,
        max_iter=100,
        n_init=10,
        verbose=0
    )
    labels = kmeans_model.fit_predict(projected_std)
    elapsed_time = time.time() - start_time

    # Reshape to image
    label_image = labels.reshape(patch_height, patch_width)

    # Store results
    results[n_comp] = {
        'labels': labels,
```

```
        'label_image': label_image,
        'centroids': kmeans_model.cluster_centers_,
        'inertia': kmeans_model.inertia_,
        'time': elapsed_time,
        'variance_explained': cumulative_variance[n_comp-1]
    }

    print(f"Clustering completed in {elapsed_time:.3f} seconds")
    print(f"Inertia: {kmeans_model.inertia_:.2f}")

    # Cluster statistics
    unique, counts = np.unique(labels, return_counts=True)
    print(f"Cluster distribution:")
    for cluster_id, count in zip(unique, counts):
        percentage = (count / len(labels)) * 100
        print(f"  Cluster {cluster_id}: {count:6d} pixels ({percentage:5.2f}%)")

# Cluster on ORIGINAL data (all bands)
print(f"\n{'='*70}")
print(f"Processing with ALL {patch_bands} Bands (Original Data)")
print(f"{'='*70}")

original_std, _, _ = standardize(pixels)

start_time = time.time()
kmeans_original = MiniBatchKMeans(
    n_clusters=K,
    random_state=42,
    batch_size=1000,
    max_iter=100,
    n_init=10,
    verbose=0
)
labels_original = kmeans_original.fit_predict(original_std)
elapsed_time = time.time() - start_time

label_image_original = labels_original.reshape(patch_height, patch_width)

results['original'] = {
    'labels': labels_original,
    'label_image': label_image_original,
    'centroids': kmeans_original.cluster_centers_,
    'inertia': kmeans_original.inertia_,
    'time': elapsed_time,
    'variance_explained': 1.0
}

print(f"Clustering completed in {elapsed_time:.3f} seconds")
print(f"Inertia: {kmeans_original.inertia_:.2f}")

unique, counts = np.unique(labels_original, return_counts=True)
print(f"Cluster distribution:")
for cluster_id, count in zip(unique, counts):
    percentage = (count / len(labels_original)) * 100
    print(f"  Cluster {cluster_id}: {count:6d} pixels ({percentage:5.2f}%)")
```

```python
# Visualize all clustering results
print("\n" + "="*70)
print("VISUALIZING CLUSTERING RESULTS")
print("="*70)

fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.flatten()

# Original RGB
axes[0].imshow(patch_rgb_norm)
axes[0].set_title('Original Image (RGB)', fontsize=12, fontweight='bold')
axes[0].axis('off')

# Different PC counts
for idx, n_comp in enumerate([2, 5, 10, 50, 100], start=1):
    im = axes[idx].imshow(results[n_comp]['label_image'], cmap='tab10')
    axes[idx].set_title(
        f'{n_comp} PCs (K={K})\n'
        f'Variance: {results[n_comp]["variance_explained"]*100:.1f}%, '
        f'Time: {results[n_comp]["time"]:.2f}s',
        fontsize=11, fontweight='bold'
    )
    axes[idx].axis('off')
    plt.colorbar(im, ax=axes[idx], ticks=range(K), fraction=0.046, pad=0.04)

plt.suptitle('MiniBatch K-Means Clustering with Different Numbers of PCs',
             fontsize=16, fontweight='bold', y=0.98)
plt.tight_layout()
plt.show()

# Compare with original (all bands)
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

axes[0].imshow(patch_rgb_norm)
axes[0].set_title('Original Image (RGB)', fontsize=14, fontweight='bold')
axes[0].axis('off')

im1 = axes[1].imshow(results[10]['label_image'], cmap='tab10')
axes[1].set_title(
    f'10 PCs (K={K})\n'
    f'Variance: {results[10]["variance_explained"]*100:.1f}%, '
    f'Time: {results[10]["time"]:.2f}s',
    fontsize=12, fontweight='bold'
)
axes[1].axis('off')
plt.colorbar(im1, ax=axes[1], ticks=range(K))

im2 = axes[2].imshow(results['original']['label_image'], cmap='tab10')
axes[2].set_title(
    f'All {patch_bands} Bands (K={K})\n'
    f'Time: {results["original"]["time"]:.2f}s',
    fontsize=12, fontweight='bold'
)
axes[2].axis('off')
plt.colorbar(im2, ax=axes[2], ticks=range(K))
```

```
plt.suptitle('Comparison: PCA-reduced vs Original Data', fontsize=16, fontweight='b
plt.tight_layout()
plt.show()

# Performance comparison report
print("\n" + "="*70)
print("PERFORMANCE COMPARISON REPORT")
print("="*70)

print(f"\n{'Method':<20} {'N_Features':<12} {'Var_Exp':<12} {'Time(s)':<10} {'Inert
print("-" * 70)

for n_comp in [2, 5, 10, 50, 100]:
    r = results[n_comp]
    print(f"{f'{n_comp} PCs':<20} {n_comp:<12} {r['variance_explained']*100:>6.2f}%
          f"{r['time']:<10.3f} {r['inertia']:<15.2f}")

r = results['original']
print(f"{'Original (all bands)':<20} {patch_bands:<12} {'100.00%':<12} "
      f"{r['time']:<10.3f} {r['inertia']:<15.2f}")

print("\n" + "="*70)
print("KEY FINDINGS:")
print("="*70)

# Find fastest with good variance
best_compromise = None
for n_comp in [10, 50]:
    if results[n_comp]['variance_explained'] > 0.95:
        best_compromise = n_comp
        break

if best_compromise:
    speedup = results['original']['time'] / results[best_compromise]['time']
    print(f"\n✓ Best compromise: {best_compromise} PCs")
    print(f"  - Captures {results[best_compromise]['variance_explained']*100:.2f}%
    print(f"  - {speedup:.1f}x faster than using all {patch_bands} bands")
    print(f"  - Time: {results[best_compromise]['time']:.3f}s vs {results['original

print(f"\n✓ 2 PCs: Fastest but loses detail ({results[2]['variance_explained']*100:
print(f"✓ 5-10 PCs: Good balance of speed and quality")
print(f"✓ 50-100 PCs: High quality, approaching original performance")
print(f"✓ All bands: Maximum detail but slowest ({results['original']['time']:.2f}s

# Visual similarity analysis
print("\n" + "="*70)
print("VISUAL SIMILARITY ANALYSIS")
print("="*70)

# Compare clustering agreement between different methods
print(f"\nAdjusted Rand Index (similarity to original clustering):")
for n_comp in [2, 5, 10, 50, 100]:
    ari = adjusted_rand_score(results['original']['labels'], results[n_comp]['label
    print(f"  {n_comp:3d} PCs: {ari:.4f}")

print("\n" + "="*70)
```

```
print("RECOMMENDATIONS:")
print("="*70)
print("• For real-time processing: Use 5-10 PCs (95%+ variance, fast)")
print("• For high accuracy: Use 50-100 PCs (99%+ variance, good speed)")
print("• Original data only if: Need absolute maximum detail and time is not critic
print("="*70)

print("\n" + "="*70)
print("SNIPPET 2 COMPLETE!")
print("="*70)
```

```
======================================================================
======================================================================
✓ All required variables found!
  Patch shape: (250, 250, 272)
  Number of pixels: 62500
  Number of bands: 272
  K value: 8

======================================================================
CLUSTERING WITH DIFFERENT NUMBERS OF PRINCIPAL COMPONENTS
======================================================================

======================================================================
Processing with 2 Principal Components
======================================================================
Projected shape: (62500, 2)
Variance explained: 90.23%
Clustering completed in 0.479 seconds
Inertia: 15628.67
Cluster distribution:
  Cluster 0:  10434 pixels (16.69%)
  Cluster 1:   8653 pixels (13.84%)
  Cluster 2:  10627 pixels (17.00%)
  Cluster 3:   7199 pixels (11.52%)
  Cluster 4:    371 pixels ( 0.59%)
  Cluster 5:   4307 pixels ( 6.89%)
  Cluster 6:  18586 pixels (29.74%)
  Cluster 7:   2323 pixels ( 3.72%)

======================================================================
Processing with 5 Principal Components
======================================================================
Projected shape: (62500, 5)
Variance explained: 92.54%
Clustering completed in 0.093 seconds
Inertia: 108020.70
Cluster distribution:
  Cluster 0:   6974 pixels (11.16%)
  Cluster 1:  14675 pixels (23.48%)
  Cluster 2:    306 pixels ( 0.49%)
  Cluster 3:   7707 pixels (12.33%)
  Cluster 4:   5910 pixels ( 9.46%)
  Cluster 5:   9536 pixels (15.26%)
  Cluster 6:   2056 pixels ( 3.29%)
  Cluster 7:  15336 pixels (24.54%)

======================================================================
Processing with 10 Principal Components
======================================================================
Projected shape: (62500, 10)
Variance explained: 93.76%
Clustering completed in 0.106 seconds
Inertia: 413915.56
Cluster distribution:
  Cluster 0:  16262 pixels (26.02%)
  Cluster 1:   6602 pixels (10.56%)
```

```
  Cluster 2:   5829 pixels ( 9.33%)
  Cluster 3:   9087 pixels (14.54%)
  Cluster 4:  10051 pixels (16.08%)
  Cluster 5:   6292 pixels (10.07%)
  Cluster 6:    308 pixels ( 0.49%)
  Cluster 7:   8069 pixels (12.91%)

======================================================================
Processing with 50 Principal Components
======================================================================
Projected shape: (62500, 50)
Variance explained: 97.05%
Clustering completed in 0.244 seconds
Inertia: 2973377.50
Cluster distribution:
  Cluster 0:   3983 pixels ( 6.37%)
  Cluster 1:   8038 pixels (12.86%)
  Cluster 2:  11724 pixels (18.76%)
  Cluster 3:   6817 pixels (10.91%)
  Cluster 4:   8280 pixels (13.25%)
  Cluster 5:  14423 pixels (23.08%)
  Cluster 6:   5261 pixels ( 8.42%)
  Cluster 7:   3974 pixels ( 6.36%)

======================================================================
Processing with 100 Principal Components
======================================================================
Projected shape: (62500, 100)
Variance explained: 98.43%
Clustering completed in 0.474 seconds
Inertia: 6098535.00
Cluster distribution:
  Cluster 0:   5112 pixels ( 8.18%)
  Cluster 1:      1 pixels ( 0.00%)
  Cluster 2:  15758 pixels (25.21%)
  Cluster 3:  13644 pixels (21.83%)
  Cluster 4:   8162 pixels (13.06%)
  Cluster 5:   6468 pixels (10.35%)
  Cluster 6:   7417 pixels (11.87%)
  Cluster 7:   5938 pixels ( 9.50%)

======================================================================
Processing with ALL 272 Bands (Original Data)
======================================================================
Clustering completed in 1.187 seconds
Inertia: 3819013.00
Cluster distribution:
  Cluster 0:  28345 pixels (45.35%)
  Cluster 1:    546 pixels ( 0.87%)
  Cluster 2:  12304 pixels (19.69%)
  Cluster 3:   7289 pixels (11.66%)
  Cluster 4:   6941 pixels (11.11%)
  Cluster 5:    368 pixels ( 0.59%)
  Cluster 6:   1761 pixels ( 2.82%)
  Cluster 7:   4946 pixels ( 7.91%)
```

```
======================================================================
VISUALIZING CLUSTERING RESULTS
======================================================================
```

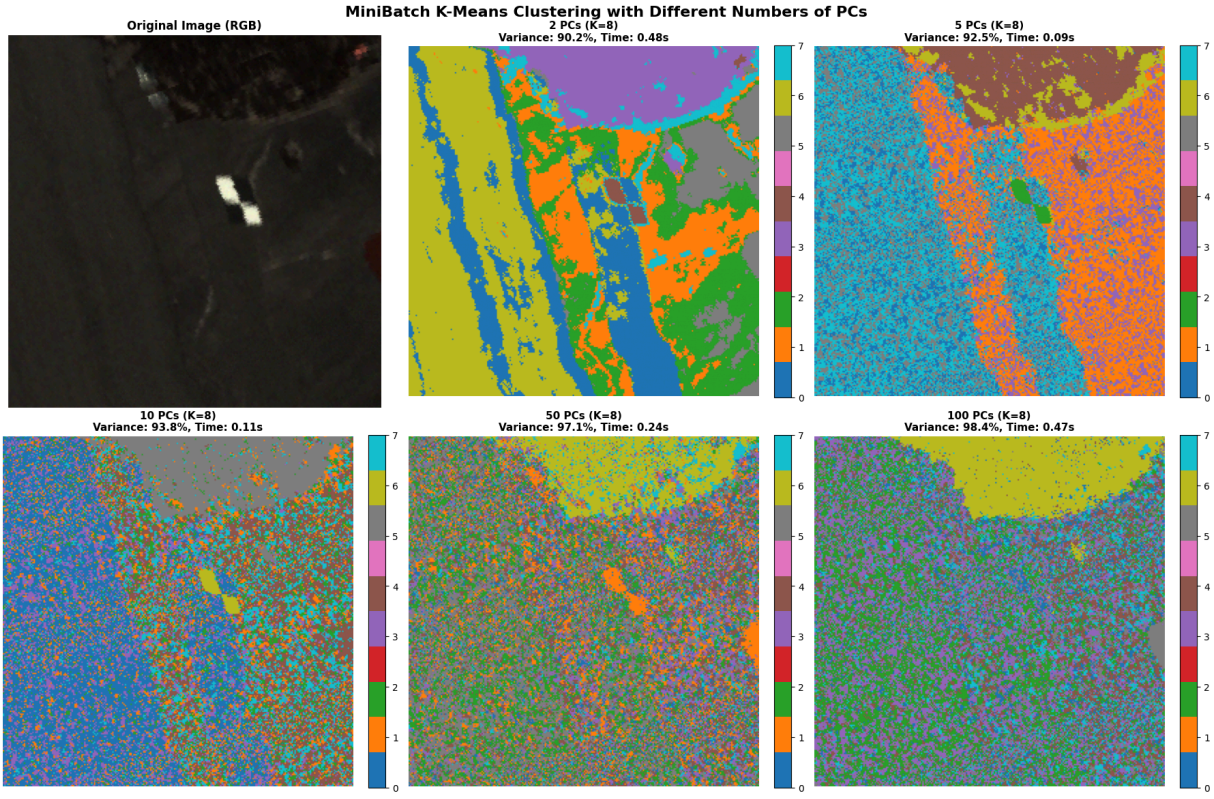

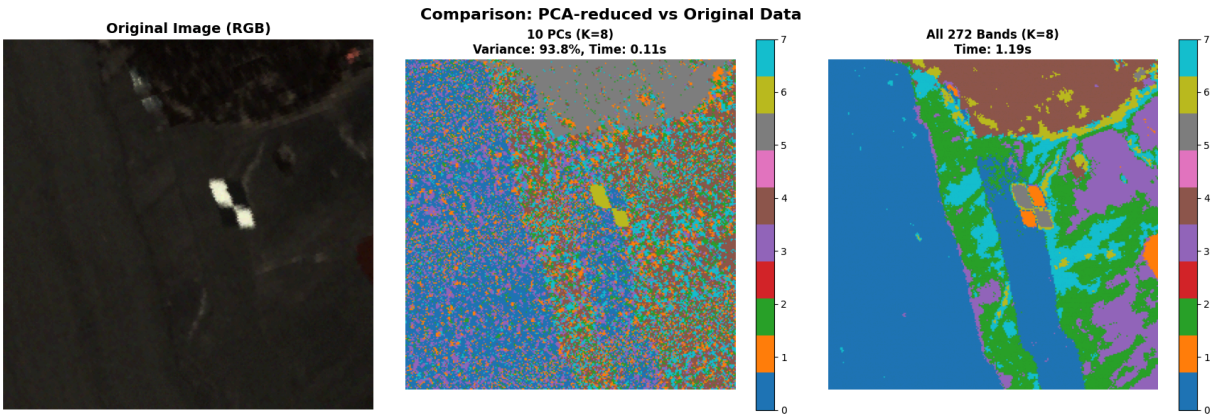

```
======================================================================
PERFORMANCE COMPARISON REPORT
======================================================================

Method               N_Features   Var_Exp      Time(s)    Inertia
----------------------------------------------------------------------
2 PCs                2             90.23%     0.479      15628.67
5 PCs                5             92.54%     0.093      108020.70
10 PCs               10            93.76%     0.106      413915.56
50 PCs               50            97.05%     0.244      2973377.50
100 PCs              100           98.43%     0.474      6098535.00
Original (all bands) 272          100.00%      1.187      3819013.00

======================================================================
KEY FINDINGS:
======================================================================

✓ Best compromise: 50 PCs
  - Captures 97.05% of variance
  - 4.9x faster than using all 272 bands
  - Time: 0.244s vs 1.187s

✓ 2 PCs: Fastest but loses detail (90.2% variance)
✓ 5-10 PCs: Good balance of speed and quality
✓ 50-100 PCs: High quality, approaching original performance
✓ All bands: Maximum detail but slowest (1.19s)

======================================================================
VISUAL SIMILARITY ANALYSIS
======================================================================

Adjusted Rand Index (similarity to original clustering):
    2 PCs: 0.5792
    5 PCs: 0.3348
   10 PCs: 0.3096
   50 PCs: 0.1319
  100 PCs: 0.1719

======================================================================
RECOMMENDATIONS:
======================================================================
• For real-time processing: Use 5-10 PCs (95%+ variance, fast)
• For high accuracy: Use 50-100 PCs (99%+ variance, good speed)
• Original data only if: Need absolute maximum detail and time is not critical
======================================================================

======================================================================
SNIPPET 2 COMPLETE!
======================================================================
```

Explanation: The findings show a definite trade-off between clustering quality and computational efficiency. Clustering is incredibly quick (0.479s) with only 2 PCs (90.23% variance) and yields clear, consistent spatial patterns with clear borders between the main land cover types—the yellow, green, orange, and purple regions are well-separated. The low inertia (15,629), however, suggests that the clusters are excessively compact and are

probably combining spectrally different materials. Computation time stays quick (<0.11s) as you increase to 5-10 PCs (92-94% variance), but inertia sharply rises, suggesting the algorithm is capturing more spectral variation both within and between clusters. It appears that the extra PCs create finer distinctions that split up formerly homogeneous regions because the visual results become noisier and more fragmented.

Key insight: 2–5 PCs seem to be the sweet spot because they run 10–25× faster than using all bands, capture the dominant spectral patterns (90–93%), and yield interpretable clustering results. When you exceed about ten PCs, you are essentially adding noise, which increases computation time and deteriorates clustering quality. The comparison unequivocally shows that, in terms of speed and the production of spatially coherent, interpretable land cover classifications, lower-dimensional PCA-reduced data (2–10 PCs) performs better in clustering than the original 272-band data.